

1 newpage

2 Data Science: Bridging Principles and Practice

2.1 Part 3: DataFrames [SOLUTIONS]

2.2 3. DataFrames

2.3 3b. Explore the data: attributes and methods

2.3.1 Attributes

EXERCISE: One extremely useful attribute is `shape`, which returns the number of rows and columns in the DataFrame, separated by commas. In the next cell, get the `shape` attribute from the `ads` DataFrame.

```
In [2]: # get the shape of the DataFrame
        data_shape = ads.shape
        data_shape
```

```
Out[2]: (588101, 6)
```

2.3.2 Methods

EXERCISE: the `describe` method is incredibly useful for learning about your data. `describe` returns summary statistics about the numerical data in your DataFrame: things like the count of non-empty items in each column, the average, the minimum, and the maximum.

In the next cell, call the `describe` method on your DataFrame. This call will look very similar to the example where the `head` method was called; only the name of the method changes.

```
In [3]: # use dot notation to call "describe" on the ads table in place of the ellipses
        ads_description = ads.describe()
        ads_description
```

```
Out[3]:
```

	user id	total ads	most ads hour
count	5.881010e+05	588101.000000	588101.000000
mean	1.310692e+06	24.820876	14.469061
...	Omitting 2 lines
50%	1.313725e+06	13.000000	14.000000
75%	1.484088e+06	27.000000	18.000000
max	1.654483e+06	2065.000000	23.000000

2.4 3c. Selecting columns

EXERCISE: Use square brackets to index the “converted” column.

```
In [4]: # index the "converted" column
        converted_col = ads["converted"]
        converted_col
```

```
Out[4]: 0          False
        1          False
        2          False
        ... Omitting 55 lines ...
        588099      False
        588100      False
        Name: converted, Length: 588101, dtype: bool
```

2.5 3d. Filtering rows

EXERCISE: Oftentimes, we want to calculate statistics separately for the control and experimental groups. Create two tables, one containing only rows where the user was in the "ad" group (the experimental group) and one with only rows where the user was in the "psa" group (the control group).

Hint: We've given you the code to create this for the experiment group. Fill in the appropriate value to select users who did NOT see ads.

```
In [5]: # users in the experiment group
        experiment = ads[ads["test group"] == "ad"]
        experiment.head()
```

```
Out[5]:
```

	user id	test group	converted	total ads	most ads day	most ads hour
0	1069124	ad	False	130	Monday	20
1	1119715	ad	False	93	Tuesday	22
2	1144181	ad	False	21	Tuesday	18
3	1435133	ad	False	355	Tuesday	10
4	1015700	ad	False	276	Friday	14

```
In [6]: # users in the control group
        # fill in the ellipses with the correct text to select users who saw psas
        control = ads[ads["test group"] == "psa"]
        control.head()
```

```
Out[6]:
```

	user id	test group	converted	total ads	most ads day	most ads hour
18	900681	psa	False	248	Saturday	19
38	905704	psa	False	27	Thursday	8
68	904595	psa	False	13	Tuesday	19
140	901904	psa	False	32	Wednesday	19
157	902234	psa	False	105	Tuesday	19